

Stan Tutorial

Module 4: Hierarchical Nested Multinomial Logit

Kevin Van Horn and Elea McDonnell Feit

June 25, 2017

Module Objectives

In this module we illustrate how Stan can be used to fit a non-standard model. This flexibility to fit almost any model you can imagine is a key advantage of Stan.

Acknowledgments



Based on work done for the Modellers along with Michael Smith.
Stan code © 2015, The Modellers, licensed under GPL v2.

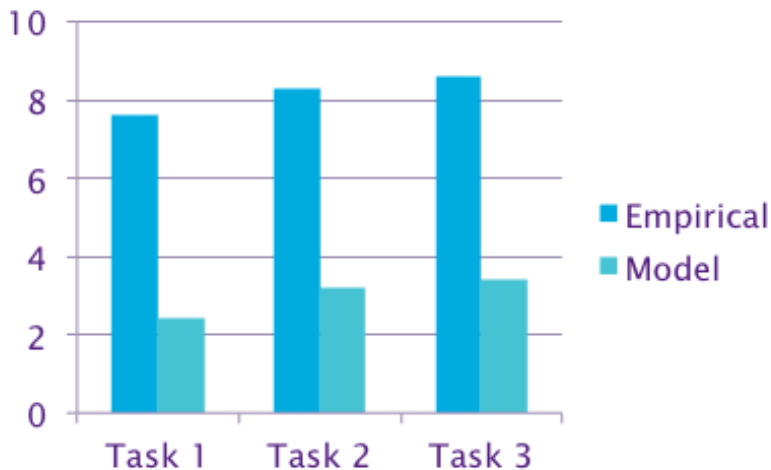


Thanks to Kevin's employer Adobe Systems for their support.

The Problem: Failed Validation Tasks

Choice-based conjoint survey; dual-response DCM.

Validation tasks: one product shown, asked if they would buy.



Posterior Predictive Probabilities

Weighted average of choice probabilities.

Weight by posterior probability of parameter vector ψ .

$$\Pr(c \mid s, i, D) = \int \Pr(c \mid s, i, \psi) p(\psi \mid D) d\psi$$

Choice c , scenario s , individual i .

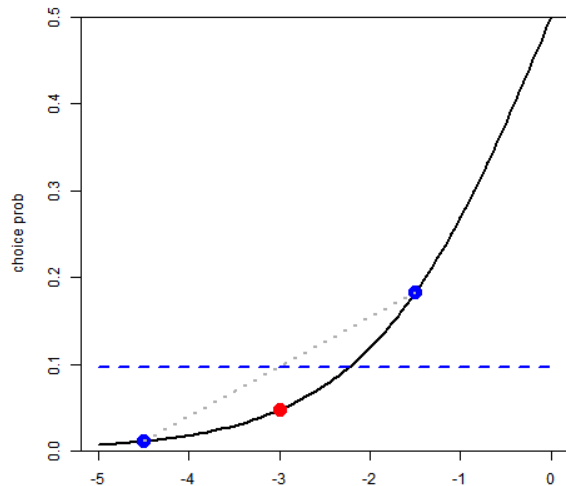
Survey data D , model parameters ψ .

Implemented by averaging predictions over posterior draws.

$$\Pr(c \mid s, i, D) \approx \frac{1}{n} \sum_{j=1}^n \Pr(c \mid s, i, \psi^{(j)})$$

Jensen's Inequality

For any convex function f , $E[f(\eta)] > f(E[\eta])$.



Example: $\frac{1}{2}f(-4.5) + \frac{1}{2}f(-1.5) > f(-3)$

Jensen's Inequality (continued)

$\Pr(\text{buy}) = f(\eta)$ where f is inverse logit, $\eta = \sum_i \beta_i x_i$

$E[f(\eta)]$ is posterior predictive choice probability.

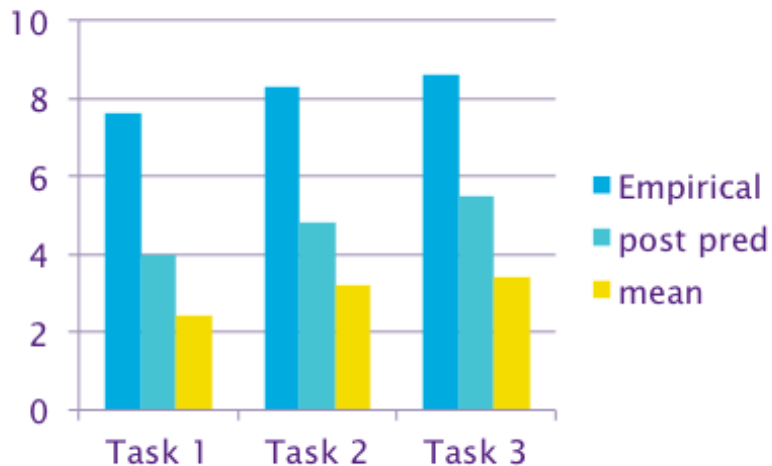
$f(E[\eta])$ is choice probability using point estimate of β .

$f(\eta)$ is convex for $\eta < 0$.

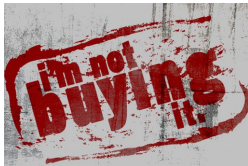
If $\eta < 0$ and using point estimate of $\beta \dots$

\dots computed choice prob $f(E[\eta])$ is *biased downward*.

Results: Using Posterior Predictive Probabilities



One of These Things Is Not Like the Others

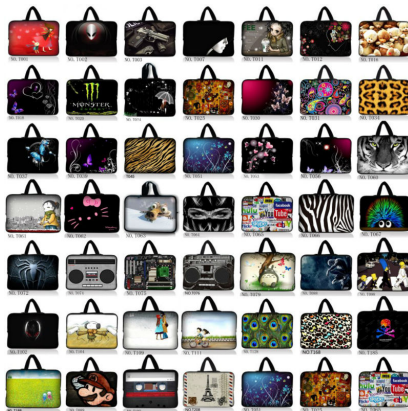


Clashes with MNL property: Independence of Irrelevant Alternatives

The Problem with Multinomial Logit

I Can Make You Buy If I Give You Enough Options

Even if they're all similar.



The Math

Let η_i be linear predictor for product i .

Linear predictor for none is 0.

$$\Pr(\text{none}) = \frac{1}{1 + \sum_{i=1}^C \exp(\eta_i)}$$

So $\Pr(\text{none}) \rightarrow 0$ as $C \rightarrow \infty$.

Tasks used for estimation: 5 product alternatives and none.

Tasks used for validation: 1 product and none.

Nested Multinomial Logit (“none” vs. all else)

$$\Pr(\text{choose product } i) = \frac{1}{1 + \exp(-\alpha - \lambda I)} \cdot \frac{\exp(\beta' \mathbf{x}_i)}{\sum_j \exp(\beta' \mathbf{x}_j)}$$

$$\Pr(\text{choose none}) = \frac{1}{1 + \exp(\alpha + \lambda I)}$$

$$I = \log \sum_j \exp(\beta' \mathbf{x}_j)$$

$0 \leq \lambda \leq 1$; if $\lambda = 1$ then MNL.

Hierarchical Nested MNL

r indexes respondents.

$$\Pr(\text{choose product } i) = \frac{1}{1 + \exp(-\alpha_r - \lambda I_r)} \cdot \frac{\exp(\beta'_r \mathbf{x}_i)}{\sum_j \exp(\beta'_r \mathbf{x}_j)}$$

$$\Pr(\text{choose none}) = \frac{1}{1 + \exp(\alpha_r + \lambda I_r)}$$

$$I_r = \log \sum_j \exp(\beta'_r \mathbf{x}_j)$$

$$\beta_{rk} \sim \text{Normal}(\boldsymbol{\theta}'_k \mathbf{z}_r, \sigma_k)$$

$$\alpha_r \sim \text{Normal}(\boldsymbol{\varphi}' \mathbf{z}_r, \sigma_\alpha)$$

$$\lambda \sim \text{Uniform}(0, 1)$$

$$\theta_{kg} \sim \text{Normal}(0, 10)$$

$$\sigma_k \sim \text{HalfNormal}(5)$$

Initial Stan Model (1)

```
data {  
  int<lower=1> R; // # respondents  
  int<lower=1> K; // # product covariates; no intercept  
  int<lower=1> G; // # respondent covariates  
  int<lower=1> S; // # scenarios per respondent  
  int<lower=2> C; // # alts (choices) per scenario  
  matrix[C, K] X[R, S];  
    // X[r,s] is cov mat of scen s for resp r.  
  matrix[G, R] Z;  
    // Z[,r] is vector of covariates for resp r  
  int<lower=1,upper=C> Y1[R,S]; // forced choice  
  int<lower=0,upper=1> Y2[R,S]; // dual response  
}
```

Initial Stan Model (2)

```
parameters {  
  real<lower=0, upper=1> lambda;  
  vector<lower=0>[K+1] sigma;  
  matrix[K+1, G] Theta;  
  matrix[K+1, R] Beta;  
}  
model {  
  lambda ~ uniform(0, 1);  
  sigma ~ normal(0, 5);  
  to_vector(Theta) ~ normal(0, 10);  
  to_vector(Beta) ~  
    normal(to_vector(Theta * Z),  
           to_vector(rep_matrix(sigma, R)));  
  ...  
}
```

Vectorization

```
to_vector(Theta) ~ normal(0, 10);  
to_vector(Beta) ~  
  normal(to_vector(Theta * Z),  
         to_vector(rep_matrix(sigma, R)));
```

is equivalent to

```
for (k in 1:(K+1))  
  for (g in 1:G)  
    Theta[k,g] ~ normal(0, 10);  
for (k in 1:(K+1))  
  for (r in 1:R)  
    Beta[k, r] ~ normal((Theta * Z)[k, r], sigma[k]);
```


Initial Stan Model (3)

```
model {  
  ...  
  for (r in 1:R) {  
    vector[K] b = Beta[2:(K+1), r];  
    real alpha = Beta[1,r];  
    for (s in 1:S) {  
      vector[C] u = X[r,s] * b;  
      real u_buy = alpha + lambda * log_sum_exp(u);  
      Y1[r,s] ~ categorical_logit(u);  
      Y2[r,s] ~ bernoulli_logit(u_buy);  
    }  
  }  
}
```

Some Stan Details

`log_sum_exp(u)` is numerically robust form of

$$\log \left(\sum_i \exp(u_i) \right)$$

`y ~ categorical_logit(u)` means

$$\Pr(y = i) = \frac{\exp(u_i)}{\sum_j \exp(u_j)}$$

`y ~ bernoulli_logit(u)` means

$$\Pr(y = 1) = \frac{\exp(u)}{1 + \exp(u)}$$

And When We Run It...

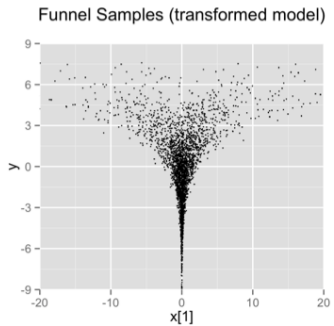
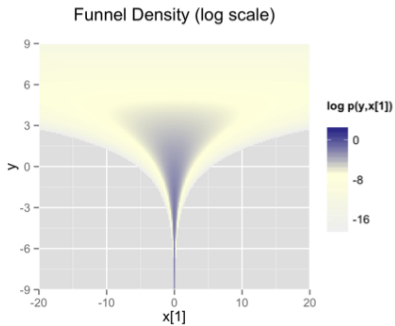
```
## There were 11 divergent transitions after warmup.  
## Increasing adapt_delta above 0.8 may help.  
## See http://mc-stan.org/misc/warnings.html#divergent-transitions  
  
## There were 4 chains where the estimated  
## Bayesian Fraction of Missing Information was low.  
## See  
## http://mc-stan.org/misc/warnings.html#bfmi-low
```

An Extreme Form of the Problem: Neal's Funnel

```
parameters {  
  real y;  
  vector[9] x;  
}  
model {  
  y ~ normal(0,3);  
  x ~ normal(0,exp(y/2));  
}
```

Variance for x strongly varies with y !

Distribution for Neal's Funnel



Why Neal's Funnel Is a Problem

Hamiltonian MC simulates trajectory of system whose potential energy is

$$V(x) = -\log(\text{prob density at } x)$$

To efficiently explore distribution, need

- ▶ large step size when curvature of $V(x)$ is small;
- ▶ small step size when curvature of $V(x)$ is large.

But have to use a single step size globally.

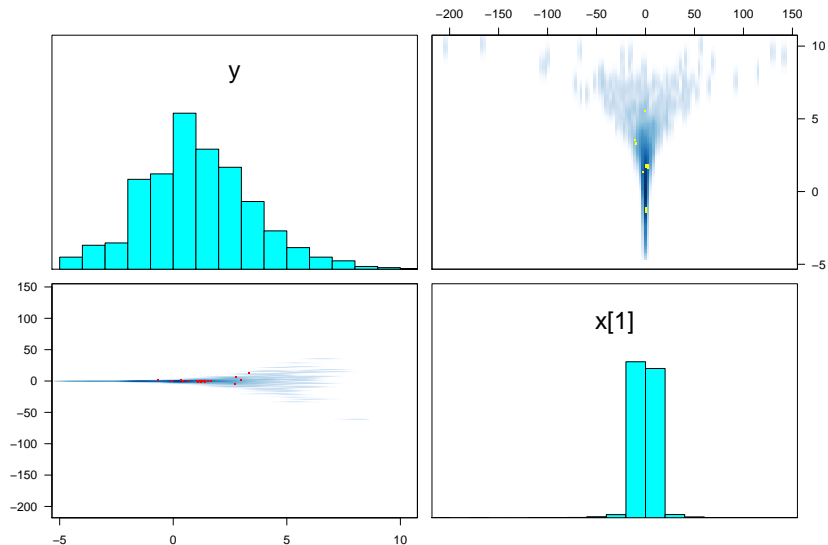
- ▶ Normal distribution: constant curvature.
- ▶ Neal's Funnel: extreme variation in curvature.

HMC has trouble entering the narrow neck of the funnel.

Warning / error message: “divergent transitions after warmup.”

Run the Funnel Model

```
fit <- stan_demo('funnel', seed=928374)
pairs(fit, pars=c('y', 'x[1]'), las=1)
```




Reparameterized Model

```
parameters {  
  real y_raw;  
  vector[9] x_raw;  
}  
transformed parameters {  
  real y = 3.0 * y_raw;  
  vector[9] x = exp(y/2) * x_raw;  
}  
model {  
  y_raw ~ normal(0,1);  
  x_raw ~ normal(0,1);  
}
```


Run the Reparameterized Model

```
fit <- stan_demo('funnel_reparam', seed=928374)  
pairs(fit, pars=c('y_raw', 'x_raw[1]'), las=1)
```



Reparameterized Stan Model (2a)

```
parameters {  
  real<lower=0, upper=1> lambda;  
  // vector<lower=0>[K+1] sigma;  
  vector<lower=0>[K+1] sigma_raw;  
  // matrix[K+1, G] Theta;  
  matrix[K+1, G] Theta_raw;  
  // matrix[K+1, R] Beta;  
  matrix[K+1, R] Epsilon;  
}
```

Reparameterized Stan Model (2b)

```
transformed parameters {  
  vector<lower=0>[K+1] sigma = 5 * sigma_raw;  
  matrix[K+1, G] Theta = 10 * Theta_raw;  
  matrix[K+1, R] Beta =  
    Theta * Z + diag_pre_multiply(sigma, Epsilon);  
}  
model {  
  lambda ~ uniform(0, 1);  
  // sigma ~ normal(0, 5);  
  sigma_raw ~ normal(0, 1);  
  // to_vector(Theta) ~ normal(0, 10);  
  to_vector(Theta_raw) ~ normal(0, 1);  
  // to_vector(Beta) ~ normal(to_vector(Theta * Z),  
  //                          to_vector(rep_matrix(sigma, R)));  
  to_vector(Epsilon) ~ normal(0, 1);  
  ...  
}
```

Results: Hnestedmnl + Posterior Predictive Probs

Histograms of $p^{(j)}$ with j running over posterior draws.

$$p^{(j)} = \frac{1}{R} \sum_{r=1}^R \Pr(\text{buy} \mid r, \psi^{(j)})$$

Each of three validation tasks.

Blue lines are empirical probs on holdout data.

